



Experiment -5

Student Name: Manjot Singh

Branch: BE-CSE

Semester: 5th

Subject Name: ADBMS

Subject Code: 23CSP-333

UID: 23BCS12549

Section/Group: KRG_2B

Date of Performance: 22/9/25

1. Problem Description/Aim:

Problem 1: Generate 1 million records per ID in 'transaction_data' using generate_series() and random(), create a normal view and a materialized view 'sales_summary' with aggregated metric, (total_quantity_sold, total_sales, total_orders), and compare their performance and execution time.

Procedure (Step-by-Step):

1. Create a large dataset:
 - Create a table names transaction_data (id, value) with 1 million records.
take id 1 and 2, and for each id, generate 1 million records in value column.
 - Use Generate_series () and random() to populate the data.
2. Create a normal view and materialized view to for sales_summary, which includes total_quantity_sold, total_sales, and total_orders with aggregation.
3. Compare the performance and execution time of both.

Sample Output Description:

The transaction_data table has 2 million rows (1 million per ID) with random values. The normal view sales_summary computes aggregates on the fly, while the materialized view sales_summary_mv stores precomputed results. Queries on the materialized view are much faster, but it needs refreshing when data changes, whereas the normal view always shows up-to-date results.

Problem 2 : Create restricted views in the sales database to provide summarized, non sensitive data to the reporting team, and control access using DCL commands(GRANT and REVOKE).

Procedure (Step-by-Step):

1. Create restricted views-
 - Define views that show only **aggregated sales data** (e.g., total_sales, total_orders) without exposing sensitive columns like customer details or payment info.
2. Assign access to reporting team(or client)-
 - Use "GRANT SELECT ON view_name TO reporting_user; " to give access.

3. Revoke access if needed.

-Use “REVOKE SELECT ON view_name FROM reporting_user;” to remove access.

4. Verify access

- Reporting users can query the view but cannot access base tables directly, ensuring security.

Sample Output Description:

The result shows the restricted view providing summarized sales data only like

- Columns shown are - product_id, total_quantity_sold, total_sales, total_orders - Columns hidden are - Customer names, addresses, payment details

A reporting user querying the view sees something like :

- Product 101 - 5000 units sold, total sales Rs. 12,50,000, 500 orders.

- Product 102 - 3200 units sold, total sales Rs. 8,60,000, 320 orders.

When the user tries to query the base “sales_transactions” table directly, access is denied, enforcing security.

2. Objective: To design and implement secure, efficient data access mechanisms by creating large-scale transaction datasets, summarizing them through normal and materialized views for performance comparison, and enforcing restricted access to sensitive data using views and DCL commands.

3. SQL QUERY AND OUTPUTS -

-- PROBLEM 1

```
Create table TRANSACTION_DATA(id int, val decimal);  
INSERT INTO TRANSACTION_DATA(ID, VAL)  
SELECT 1, RANDOM()  
FROM GENERATE_SERIES(1, 1000000);
```

```
INSERT INTO TRANSACTION_DATA(ID, VAL)  
SELECT 2, RANDOM()  
FROM GENERATE_SERIES(1, 1000000);  
SELECT * FROM TRANSACTION_DATA;  
CREATE or REPLACE VIEW SALES_SUMMARY AS  
SELECT  
ID,
```



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
COUNT(*) AS total_quantity_sold,  
sum(val) AS total_sales,  
count(distinct id) AS total_orders  
FROM TRANSACTION_DATA GROUP  
BY ID;
```

```
EXPLAIN ANALYZE  
SELECT * FROM SALES_SUMMARY;
```

```
CREATE MATERIALIZED VIEW SALES_SUMM AS  
SELECT  
ID,  
COUNT(*) AS total_quantity_sold,  
sum(val) AS total_sales,  
count(distinct id) AS total_orders  
FROM TRANSACTION_DATA GROUP  
BY ID;
```

```
EXPLAIN ANALYZE  
SELECT * FROM SALES_SUMM;
```

-- PROBLEM 2

```
CREATE TABLE customer_data (  
    transaction_id SERIAL PRIMARY KEY,  
    customer_name VARCHAR(100),  
    email VARCHAR(100),  
    phone VARCHAR(15),  
    payment_info VARCHAR(50),  
    order_value DECIMAL,  
    order_date DATE DEFAULT CURRENT_DATE  
);
```

```
INSERT INTO customer_data (customer_name, email, phone, payment_info, order_value)  
VALUES  
( 'M', 'M@example.com', '9131094977', '1234-5678-9012-3456', 500),  
( 'A', 'A@example.com', '9931094977', '1234-5678-9012-3456', 234),
```



DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
('B', 'B@example.com', '9263444151', '9876-5432-1098-7654', 754),  
('C', 'C@example.com', '9263444151', '9876-5432-1098-7654', 300);
```

```
CREATE OR REPLACE VIEW restricted_sales_data AS  
SELECT  
    customer_name,  
    COUNT(*) AS total_orders,  
    SUM(order_value) AS total_sales  
FROM customer_data  
GROUP BY customer_name;
```

```
SELECT * FROM restricted_sales_data;
```

```
CREATE USER client1 WITH PASSWORD 'password123';
```

```
GRANT SELECT ON restricted_sales_data TO client1;  
REVOKE SELECT ON restricted_sales_data FROM client1;
```

Output:

Output 1:

Output:

```
CREATE TABLE  
INSERT 0 1000  
INSERT 0 1000
```

id	val
1	0.85769666100434
1	0.156152411912014
1	0.769247939305863
1	0.94759144873942
1	0.535303978648415
1	0.00759400169134139
1	0.99270863754365
1	0.239119714886034

Output 2:

Output:

```
CREATE TABLE
```

```
INSERT 0 4
```

```
CREATE VIEW
```

customer_name	total_orders	total_sales
B	1	754
C	1	300
M	1	500
A	1	234

(4 rows)

```
psql:commands.sql:28: ERROR:  permission denied to create role
```

Learning Outcomes:

- Successfully implemented sub-queries to extract top salary earners by department.
- Successfully implemented stored procedures in PostgreSQL.
- Practiced handling input and output parameters in procedures.
- Automated HR analytics queries for gender-based employee counts.
- Developed an order-processing system with real-time stock validation.
- Enhanced SQL procedural programming skills for enterprise applications.